

# Stephanie Milani

✉ [smilani@andrew.cmu.edu](mailto:smilani@andrew.cmu.edu)  
📄 <https://stephmilani.github.io/>  
🌐 [stephmilani](#)

## Research Interests . . . . .

Reinforcement learning, interpretability and explainability, AI evaluation and measurement, AI transparency, human-AI interaction, sequential decision-making.

## Education . . . . .

2019–Present **Carnegie Mellon University.**

Ph.D. Candidate in Machine Learning.

Advisor: Fei Fang.

Committee: Fei Fang, Hong Shen, Geoffrey J. Gordon, Katja Hofmann, & Oriol Vinyals.

2021 **Carnegie Mellon University.**

M.S. in Machine Learning Research.

2019 **University of Maryland, Baltimore County.**

B.S. in Computer Science, B.A. in Psychology. *Cum Laude*. Honors College Certificate. Presidential Scholar.

Advisors: Marie desJardins and Cynthia Matuszek.

## Research Experience . . . . .

2019–Present **Carnegie Mellon University, Machine Learning Department.**

Graduate Researcher advised by Fei Fang.

2022 **Microsoft Research, Montreal, Reinforcement Learning Group.**

Ph.D. Research Intern advised by Geoffrey J. Gordon.

2021 **Microsoft Research, Cambridge, Deep Reinforcement Learning for Games Group.**

Ph.D. Research Intern advised by Katja Hofmann and Harm van Seijen.

2019 **Carnegie Mellon University, Robotics Institute.**

Post-Bacc Research Intern advised by David Held.

2018–2019 **Carnegie Mellon University, Robotics Institute.**

Undergraduate Research Intern advised by Katia Sycara, *Robotics Institute Summer Scholars Program*. Supported by NSF REU.

**University of Maryland, Baltimore County, Computer Science and Electrical Engineering.**

Undergraduate Research Assistant advised by Cynthia Matuszek. *URA Scholar*.

2016–2019 **University of Maryland, Baltimore County, Computer Science and Electrical Engineering.**

Undergraduate Research Assistant advised by Marie desJardins.

2017 **Carnegie Mellon University, Robotics Institute.**

Undergraduate Research Intern advised by Christoph Mertz, *Robotics Institute Summer Scholars Program*.

2014–2016 **University of Maryland, School of Medicine, Neurobiology Department.**

Undergraduate Research Assistant advised by Jennifer Wenzel.

## Publications . . . . .

\* denotes equal contribution.

### Book Chapters

- [1] **S. Milani**, Z. Zhang, N. Topin, Z. R. Shi, C. Kamhoua, E. E. Papalexakis, F. Fang. Interpretable Multi-Agent Reinforcement Learning with Decision-Tree Policies. *Explainable Agency in Artificial Intelligence*, CRC Press / Taylor & Francis, 2024.

### Refereed Publications

- [1] M. Phan, K. Brantley, **S. Milani**, S. Mehri, G. Swamy, G. Gordon. When is Transfer Learning Possible? *International*

*Conference on Machine Learning (ICML)*, 2024. 27.5% acceptance.

- [2] A. Venugopal, **S. Milani**, F. Fang, B. S. Ravindran. MABL: Bi-Level Latent-Variable World Model for Sample-Efficient Multi-Agent Reinforcement Learning. *International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)*, 2024. 25% acceptance.
- [3] **S. Milani**, N. Topin, M. Veloso, F. Fang. Explainable Reinforcement Learning: A Survey and Comparative Review. *ACM CSUR Special Issue on Trustworthy AI*, 2024.
- [4] **S. Milani**, A. Kanervisto, K. Ramanauskas, S. Schulhoff, B. Houghton, R. Shah. The MineRL BASALT Evaluation and Demonstrations Dataset for Training and Benchmarking Agents that Solve Fuzzy Tasks. *Neural Information Processing Systems (NeurIPS) Benchmarks & Datasets Track*, 2023. **Oral Presentation**, 1% acceptance.
- [5] **S. Milani**, A. Juliani, I. Momennejad, R. Georgescu, J. Rzepcki, A. Shaw, G. Costello, F. Fang, S. Devlin, K. Hofmann. Navigates Like Me: Understanding How People Evaluate Human-Like AI in Video Games. *ACM Conference on Human Factors in Computing Systems (CHI)*, 2023. 27.7% acceptance.  
*Previous versions in NeurIPS Workshop on Human-Centered AI, 2021; CHI Late Breaking Work, 2022.*
- [6] M. Carroll, O. Paradise, J. Lin, R. Georgescu, M. Sun, D. Bignell, **S. Milani**, K. Hofmann, M. Hausknecht, A. Dragan, S. Devlin. Uni[MASK]: Unified Inference in Sequential Decision Problems. *Neural Information Processing Systems (NeurIPS)*, 2022. **Oral Presentation**, 1.9% acceptance.  
*Previous version in ICLR Workshop on Generalizable Policy Learning in the Physical World, 2022.*
- [7] Y. Du, Z. Song, **S. Milani**, C. Gonzalez, F. Fang. Learning to Play Adaptive Cyber Deception Game. *The 13th Workshop on Optimization and Learning in Multiagent Systems at the International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2022.
- [8] **S. Milani\***, Z. Zhang\*, N. Topin, Z. R. Shi, C. Kamhoua, E. Papalexakis, F. Fang. MAVIPER: Learning Decision Tree Policies for Interpretable Multi-Agent Reinforcement Learning. *European Conference on Machine Learning (ECML)*, 2022. 26% acceptance.  
*Previous version in AAAI Explainable Agency in AI Workshop, 2022, Invited for book chapter.*
- [9] N. Topin, **S. Milani**, F. Fang, M. Veloso. Iterative Bounding MDPs: Learning Interpretable Policies via Non-Interpretable Methods. *AAAI Conference on Artificial Intelligence*, 2021. 21% acceptance.
- [10] **S. Milani\***, Z. Fan\*, S. Gulati, T. Nguyen, F. Fang, A. Yadav. Intelligent Tutoring Strategies for Students with Autism Spectrum Disorder: A Reinforcement Learning Approach. *AAAI Conference on Artificial Intelligence (AAAI) Workshop on AI for Education*, 2020.  
*Also accepted for lightning talk at the CMU Symposium on AI and Social Good, 2020.*
- [11] **S. Milani**, W. Shen, K. S. Chan, S. Venkatesan, N. O. Leslie, C. Kamhoua, F. Fang. Harnessing the Power of Deception in Attack Graph Games. *Conference on Decision and Game Theory for Security (GameSec)*, 2020.
- [12] J. Winder, **S. Milani**, M. Landen, E. Oh, S. Parr, S. Squire, M. desJardins, C. Matuszek. Planning with Abstract Learned Models While Learning Transferable Subtasks. *AAAI Conference on Artificial Intelligence*, 2020. 20.6% acceptance.  
*Previous versions in ICAPS IntEx Workshop, 2017; RLDM, 2017; and Do Good Robotics Symposium, 2019.*
- [13] B. Houghton, **S. Milani**, N. Topin, W. H. Guss, K. Hofmann, D. Perez-Liebana, M. Veloso, R. Salakhutdinov. Guaranteeing Reproducibility in Deep Learning Competitions. *Neural Information Processing Systems (NeurIPS) Challenges in Machine Learning (CiML) Workshop*, 2019.
- [14] **S. Milani**, N. Topin, K. Sycara. Penalty-Modified Markov Decision Processes: Efficient Incorporation of Norms into Sequential Decision Making Problems. *Multidisciplinary Conference on Reinforcement Learning and Decision Making (RLDM)*, 2019.
- [15] H. Li, **S. Milani**, V. Krishnamoorthy, M. Lewis, K. Sycara. Perceptions of Domestic Robots' Normative Behavior Across Cultures. *AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES)*, 2019.

#### Refereed Competition Papers and Proceedings

- [1] **S. Milani**, A. Kanervisto\*, K. Ramanauskas\*, S. Schulhoff, B. Houghton, S. Mohanty, B. Galbraith, K. Chen, Y. Song, T. Zhou, B. Yu, H. Liu, K. Guan, Y. Hu, T. Lv, F. Malato, F. Leopold, A. Raut, V. Hautamäki, A. Melnik,

- S. Ishida, J. F. Henriques, R. Klassert, W. Laurito, E. Novoseller, V. G. Goecks, N. Waytowich, D. Watkins, J. Miller, R. Shah. Towards Solving Fuzzy Tasks with Human Feedback: A Retrospective of the MineRL BASALT 2022 Competition. To appear in *PMLR: NeurIPS 2022 Competition and Demonstration Track*, 2024.
- [2] A. Kanervisto, **S. Milani**, K. Ramanauskas, B. V. Galbraith, S. H. Wang, B. Houghton, S. Mohanty, R. Shah. The MineRL BASALT Competition on Learning from Human Feedback. *Conference on Neural Information Processing Systems (NeurIPS) Competition Track*, 2022.
- [3] R. Shah, S. H. Wang, C. Wild, **S. Milani**, A. Kanervisto, V. G. Goecks, N. Waytowich, D. Watkins-Valls, B. Prakash, E. Mills, D. Garg, A. F. Fries, A. Souly, J. S. Chan, D. del Castillo, T. Lieberum. Retrospective on the 2021 BASALT Competition on Learning from Human Feedback. *PMLR: NeurIPS 2021 Competition and Demonstration Track*, 2022.
- [4] A. Kanervisto\*, **S. Milani\***, K. Ramanauskas, N. Topin, Z. Lin, J. Li, J. Shi, D. Ye, Q. Fu, W. Yang, W. Hong, Z. Huang, H. Chen, G. Zeng, Y. Lin, V. Micheli, E. Alonso, F. Fleuret, A. Nikulin, Y. Belousov, O. Svidchenko, A. Shpilman. MineRL Diamond 2021 Competition: Overview, Results, and Lessons Learned. *PMLR: NeurIPS 2021 Competition and Demonstration Track*, 2022.
- [5] W. H. Guss, A. Dirik\*, B. V. Galbraith\*, B. Houghton\*, A. Kanervisto\*, N. S. Kuno\*, **S. Milani\***, S. Mohanty\*, K. Ramanauskas\*, R. Salakhutdinov\*, R. Shah\*, N. Topin\*, S. H. Wang\*, C. Wild\*. The MineRL Diamond Competition on Sample Efficient Reinforcement Learning. *Conference on Neural Information Processing Systems (NeurIPS) Competition Track*, 2021.
- [6] R. Shah, C. Wild, S. H. Wang, N. Alex, B. Houghton, W. H. Guss, **S. Milani**, N. Topin, P. Abbeel, S. Russell, A. Dragan. The MineRL BASALT Competition on Learning from Human Feedback. *Conference on Neural Information Processing Systems (NeurIPS) Competition Track*, 2021.
- [7] W. H. Guss, **S. Milani**, N. Topin, B. Houghton, S. Mohanty, A. Melnik, A. Harter, B. Buschmaas, B. Jaster, C. Berganski, D. Heitkamp, M. Henning, H. Ritter, C. Wu, X. Hao, Y. Lu, H. Mao, Y. Mao, C. Wang, M. Opanowicz, A. Kanervisto, Y. Schraner, C. Scheller, X. Zhou, L. Liu, D. Nishio, T. Tsuneda, K. Ramanauskas, G. Juceviciute. Towards robust and domain agnostic reinforcement learning competitions: MineRL 2020. *PMLR: NeurIPS 2020 Competition and Demonstration Track*, 2021.
- [8] W. H. Guss, S. Devlin\*, B. Houghton\*, N. S. Kuno\*, **S. Milani\***, S. Mohanty\*, R. Salakhutdinov\*, J. Schulman\*, N. Topin\*, O. Vinyals\*. NeurIPS 2020 Competition: The MineRL Competition on Sample Efficient Reinforcement Learning using Human Priors. *Conference on Neural Information Processing Systems (NeurIPS) Competition Track*, 2020.
- [9] **S. Milani**, N. Topin, B. Houghton, W. H. Guss, S. P. Mohanty, O. Vinyals, N. S. Kuno. A Retrospective Analysis of the 2019 MineRL Competition on Sample-Efficient Reinforcement Learning Using Human Priors. *PMLR: NeurIPS 2019 Competition and Demonstration Track*, 2020.
- [10] W. H. Guss, C. Codel\*, K. Hofmann\*, B. Houghton\*, N. S. Kuno\*, **S. Milani\***, S. Mohanty\*, D. Perez-Liebana\*, R. Salakhutdinov\*, N. Topin\*, M. Veloso\*, and P. Wang\*. The MineRL Competition on Sample Efficient Reinforcement Learning using Human Priors. *Conference on Neural Information Processing Systems (NeurIPS) Competition Track*, 2019.

### Technical Reports

- [1] **S. Milani**. Penalty-Modified Abstract Markov Decision Processes. *Technical Report for Undergraduate Research Award, UMBC*, 2019.
- [2] **S. Milani**. Creating a Scalable Framework for Model-Free Reinforcement Learning in Norm-Rich Environments. *Robotics Institute Summer Scholars Working Papers Journal*, 2018.
- [3] **S. Milani** and C. Mertz. Generating Hard Positive Examples via Adversary for Occluded Traffic Sign Detection. *Robotics Institute Summer Scholars Working Papers Journal*, 2017.

### Honors and Awards . . . . .

#### Research

2024 **Future Leader in Responsible Data Science & AI**, UMichigan Institute for Data Science.  
*Selected as 1 of around 40 outstanding data science and AI researchers.*

- 2023 **NeurIPS Scholar Award**  
**Graduate Student Travel Grant**, Carnegie Mellon University.  
*\$750 USD to attend CHI 2023.*
- 2022 **ECML PhD Registration Grant**, 17% acceptance.  
**AAAI Student Scholarship**
- 2019 **RLDM Travel Award**.  
**ICML Travel Award**.  
**NeurIPS Travel Award**.  
**Top Reviewer Award**, ICML.  
*Awarded to top 33% of all reviewers.*
- 2018–2019 **Undergraduate Research Award**, University of Maryland, Baltimore County.  
*Selective \$1500 USD scholarship to support original undergraduate research with a faculty mentor. 1 of 55 total awarded.*
- 2018 **Inclusion@RSS Scholar**  
**Oracle Scholar**, OurCS Workshop.  
*1 of 102 chosen nationally to attend OurCS research workshop. Of those, one of a select group of participants designated as Oracle Scholar.*  
**Researcher of the Week**, University of Maryland, Baltimore County.
- 2017 **Researcher of the Week**, University of Maryland, Baltimore County.  
**Women in Transportation Fellow**, Traffic 21.  
*1 of 1 awarded to support RISS internship.*  
 Academic and Service
- 2021 **Funding Proposal for MineRL BASALT**, AIJ.  
*Awarded to R. Shah, C. Wild, S. H. Wang, N. Alex, B. Houghton, W. H. Guss, S. Milani, N. Topin, P. Abbeel, S. Russell, A. Dragan.*
- 2020 **Inclusivity Compute and Conference Grants for MineRL Competition at NeurIPS**, AIJ.  
*Awarded €7,500 to S. Milani, N. Topin.*
- 2019 **Inclusivity Travel Grants for MineRL Competition and Workshop at NeurIPS**, AIJ.  
*Awarded €3,000 to S. Milani, N. Topin.*  
**Community Service Award**, University of Maryland, Baltimore County Honors College.  
*1 of 4 awarded at UMBC for strong academic performance and outstanding community service.*  
**Grace Hopper Student Scholar**, Grace Hopper Conference.  
*Funded by Palo Alto Networks to attend Grace Hopper Conference. 1 of 657 awarded nationally.*
- 2018–2019 **Rewriting the Code Fellow**  
**Newman Civic Fellow**, Campus Compact.  
*1 of 268 awarded nationally for leadership and dedication to increasing access to computer science. Nominated by the President of UMBC, Dr. Freeman Hrabowski.*
- 2017–2019 **NAE Grand Challenge Scholar**, University of Maryland, Baltimore County.  
*1 of 7 students selected in yearly cohort to focus on finding solutions to address important societal problems.*
- 2017–2018 **France-Merrick Scholar**, UMBC.  
*1 of 7 awarded. For commitment to leadership and service in computer science and artificial intelligence.*

---

## Professional and Community Service . . . . .

### Organizing

- 2023 **WiML Un-Workshop**, *Logistics Chair*, ICML.  
**WiML Un-Workshop**, *Co-Organizer of Evening Social*, ICML.  
*With A. Bunnell.*

- 2022 **MineRL BASALT Competition on Fine-Tuning from Human Feedback**, *Workshop Co-Organizer and Panel Moderator*, NeurIPS.  
*Panelists: Kianté Brantley, Sam Devlin, Fei Fang, Oriol Vinyals.*
- 2021–2022 **MineRL BASALT Competition on Learning from Human Feedback**, *Competition Organizer*, NeurIPS.  
*2022: With A. Kanervisto, K. Ramanauskas, B. V. Galbraith, S. H. Wang, S. Schulhoff, B. Houghton, S. P. Mohanty, R. Shah.*  
*2021: With R. Shah, C. Wild, S. H. Wang, N. Alex, B. Houghton, W. H. Guss, S. P. Mohanty, A. Kanervisto, N. Topin, P. Abbeel, S. Russell, A. Dragan.*
- 2019–2021 **MineRL Competition on Sample-Efficient Reinforcement Learning**, *Competition Organizer*.
- 2021 **Minecraft as a Research Platform for RL**, *Workshop Organizer*, MSR Research Summit.  
*With E. Zuniga.*
- 2017–2018 **Meet Your Professor Speaker Series**, *Seminar Organizer*. UMBC.
- Reviewing  
**AAMAS Workshop on Reinforcement Learning in Games**, 2024.  
**ICML**, 2024, 2023, 2022, 2021, 2020.  
**NeurIPS**, 2023.  
**AAAI PC Member of Student Program**, 2023.  
**GameSec**, 2022.  
**AAMAS Workshop on Autonomous Agents for Social Good**, 2022.  
**Cooperative AI Foundation Proposal**, 2021.  
**NeurIPS Competition Track**, 2021.  
**ICLR**, 2020.  
**AAAI**, 2020.  
**Game Theory and Machine Learning for Cyber Security (Book Chapter Reviewer)**, 2020.  
**AAAI Workshop on Diversity in Artificial Intelligence**, 2020.  
**ICML Workshop on AI for Social Good**, 2019.  
**ICLR Workshop on AI for Social Good**, 2019.  
**RISS Working Papers Journal**, 2018, 2017.
- Academic and Departmental Service
- 2023–2024 **Carnegie Mellon University Machine Learning PhD Admissions Committee**.
- 2022–2023 **Carnegie Mellon University Machine Learning PhD Admissions Committee**.
- 2021–2022 **Carnegie Mellon University Machine Learning PhD Admissions Committee**.
- 2020–2021 **Carnegie Mellon University Machine Learning Master's Admissions Committee**.
- 2020 **Carnegie Mellon University RISS Admissions Committee**.
- 2018 **RISS Working Papers Journal**, Assistant Managing Editor. CMU.
- 2017 **RISS Working Papers Journal**, Assistant Managing Editor. CMU.
- Outreach
- 2023, 2020 **CMU Graduate Application Support Program (GASP)**, Mentor.  
*Supported student now PhD student at the University of Cambridge.*
- 2019–2021 **CMU AI Mentorship Program**, Mentor.
- 2019 **Steel City Showdown FIRST Robotics Competition**, Referee and Volunteer.  
**Rewriting the Code Alumni Office Hours**, Mentor.  
*Dedicated 30 min/week to provide career and academic advice to female undergrads.*
- 2016–2019 **UMBC Computer Science Education**, Vice President (2016–2017), President (2017–2018), Treasurer (2018–2019).
- 2017 **Creative Coders**, Co-founder. Co-founded program to introduce middle-school students to CS.

**QuHacks Hackathon at UMBC, Organizer.**

*Organized day-long hackathon for appx. 100 high-school and middle-school students.*

**North County High School Computer Science Classes, Co-presenter.**

*Presented on various CS topics to high school students.*

2017, 2016 **Hour of Code at UMBC, Organizer and Volunteer.**

*Organized and volunteered during two-day-long Hour of Code events on CS and AI.*

---

## Invited Talks and Events. . . . .

Does not include contributed conference or workshop talks.

2024 **Guest Lecture**, 17-537/17-737 AI Methods for Social Good at Carnegie Mellon University.

*Explainable AI for Social Good.*

2023 **Invited Speaker**, Miao Embodied AI Lab at the University of Connecticut.

*Human-Interpretable Reinforcement Learning.*

**Invited Speaker**, Berkeley Multi-Agent Learning Seminar.

*Human-Interpretable Multi-Agent Reinforcement Learning.*

**Invited Speaker**, AMD.

*Navigates Like Me: Understanding How People Evaluate Human-Like AI in Video Games.*

**Invited Speaker**, University of Maryland Multi-Agent Reinforcement Learning Seminar.

*Human-Interpretable Multi-Agent Reinforcement Learning.*

**Selected Hot Desk Demo**, CHI Interactivity.

*Navigates Like Me: Understanding How People Evaluate Human-Like AI in Video Games.*

2022 **Invited Speaker**, Women in AI Ignite at NeurIPS.

*Human-Centered Multi-Agent Systems: Learning from and for People.*

**Invited Panelist**, Machine Learning Graduate Student Panel, Carnegie Mellon University Robotics Institute Summer Scholars Program.

2021 **Invited Speaker**, Microsoft Research AI and Gaming Research Summit.

*The MineRL 2020 Competition on Sample Efficient Reinforcement Learning using Human Priors.*

**Invited Panelist**, Microsoft Research AI and Gaming Research Summit.

**Invited Panelist**, Graduate School Application Support, Carnegie Mellon University Robotics Institute Summer Scholars Program.

2020 **Invited Participant**, SOCML.

*Doing 'cognitive neuroscience' on models - will it help us understand generalization?.*

**Invited Participant**, SOCML.

*Deep Reinforcement Learning.*

**Invited Speaker**, The Campus Laboratory School at Carlow University Career Day.

*Sustainability through Computer Science.*

2019 **Invited Panelist**, RISS Orientation, Carnegie Mellon University Robotics Institute Summer Scholars Program.

2018 **Invited Participant**, CCC AI Roadmap Workshop: Integrated Intelligence. Sole undergraduate student. Resulted in *A 20-Year Community Roadmap for AI Research in the US.*

---

## Teaching Experience and Preparation . . . . .

2024–Present **Eberly Future Faculty Program, CMU, Participant.**

2024–Present **Machine Learning with Large Datasets (10-405/10-605), CMU, Co-Head Teaching Assistant.**  
Instructors: Geoff Gordon and Ameet Talwalkar.

2021 **Historical Advances in Machine Learning (10-777), CMU, Teaching Assistant.**

Instructor: Aaditya Ramdas.

2017–2018 **Creative Coders, Curriculum Developer.**

Developed curriculum for middle-school students to learn CS concepts.  
2016–2017 **Computer Science Matters in Maryland**, Curriculum Developer.

■ Affiliations . . . . .

SIGCHI, Rewriting the Code Alumni, AAAI, ACM.

■ Selected Media Coverage . . . . .

- “How to design artificial intelligence that acts nice — and only nice,” by Kathryn Hulick. Science News Explores. 2024.
- “Security games reveal how networks can fool cyber attackers.” U.S. Army DEVCOM Army Research Laboratory Public Affairs. 2020.
- “MineRL sample-efficient reinforcement learning challenge—back for a second year—benefits organizers, as well as larger research community,” by Noboru Sean Kuno. Microsoft Research Blog. 2020.
- “AI takes on popular Minecraft game in machine-learning contest,” by Jeremy Hsu. Nature Journal. 2019.
- “Project Malmo competition returns with student organizers and a new mission: To democratize reinforcement learning,” by Noboru Sean Kuno. Microsoft Research Blog. 2019.
- “Traffic21's Women in Transportation Awardee Joining CMU's Machine Learning Ph.D. Program.” Mobility21. 2019.
- “Stephanie Milani named Newman Civic Fellow for expanding access to CS education,” by Catalina Sofia Dansberger Duque. UMBC News. 2018.
- “The Hour of Code Arrives at UMBC,” by Declan Keefe. The Retriever. 2017.